

# Congestion Control

# Module Goals


At the conclusion of this module, students will be able to

- ▶ define congestion and explain the primary source of congestion
- ▶ calculate the fairness of a set of network flows

# A Sample Network

- ▶ consider two hosts communicating with two other hosts through a router with infinite buffers
  - ▶ router can transmit at  $R$  bps
  - ▶ hosts can transmit at  $T$  bps
- ▶ three possible cases:

$$T \ll R/2 \quad T \gg R/2 \quad T \approx R/2$$



# Congestion Consequences

- ▶ what happens in each situation?
- ▶  $T \ll R/2$ : if the hosts are sending much less than the output rate of the router, things will be okay
- ▶  $T \gg R/2$ : if the hosts are sending much more than the output rate, things will be really bad
- ▶  $T \approx R/2$ : if the hosts are sending just about as fast as the router can output, things will be pretty bad too

**How can we model this?**

## A Quick Aside: M/M/1 Queues

- ▶ model of a queue!  
(typically assuming the queue is infinitely sized!)
- ▶ time between arrivals follows an exponential distribution:

$$f_{\lambda} = \lambda e^{-\lambda t}$$

with average time  $1/\lambda$

- ▶ time to service an arrival follows an exponential distribution:

$$g_{\mu} = \mu e^{-\mu t}$$

with average time  $1/\mu$

# Congestion Consequences

- ▶ the average number of items in the queue is

$$L_q = \rho^2 / (1 - \rho)$$

where  $\rho = \lambda / \mu$


- ▶ if the hosts are sending much less than the output rate of the router, things will be okay  
(if  $\lambda \ll \mu$ , the queue services arrivals quickly)
- ▶ surprisingly, if the hosts are sending just about as fast as the router can output, things are pretty bad too!  
( $\lambda \rightarrow \mu$ , i.e.,  $\rho \rightarrow 1$ , the average number of items explodes towards infinity!)
- ▶ congestion occurs as packet arrival rate nears link capacity

# Costs of Congestion

- ▶ consider the same network, but now the router has finite buffers
- ▶ as the buffers fill, segments get dropped
- ▶ hosts will retransmit dropped segments when discovered
- ▶ every time we retransmit something, we are not transmitting something new
- ▶ if we only transmit dropped segments, our transmission rate of new segments drops to about  $R/3$
- ▶ if we erroneously retransmit non-dropped segments, the effective transmission rate drops to about  $R/4$

# Costs of Congestion

- ▶ consider a more complicated network...
- ▶ we can starve the A–C connection if the B–D traffic is greater than the output from R1
- ▶ if packets from the A–C connection are constantly dropped, the work done by R1 is pointless and upstream capacity is wasted





# Our Goal

- ▶ if congestion is the problem, congestion control becomes our goal
- ▶ to do that, we need to effectively use the network resources available to us:
  - ▶ the bandwidth of the links
  - ▶ the buffer space at the routers

# The Players

- ▶ doing this requires effort from both the network and the hosts!
- ▶ network responsibilities: various queuing disciplines can be used to control the order in which packets get transmitted and which packets get dropped
- ▶ host responsibilities: use congestion control mechanisms to pace how fast sources are allowed to send packets

# A Flow-Oriented View

- ▶ we're going to assume that the network is essentially connectionless, with any connection-oriented service implemented in the transport protocol
- ▶ however, this requires us to update our definition of connectionless a little
- ▶ classically, all datagrams are completely independent in a connectionless network  
(datagrams are switched independently, but who sends just one datagram?)
- ▶ there is realistically a **flow** of related datagrams through the network

# A Flow-Oriented View

- ▶ flows can be abstracted at different granularity
  - ▶ host-to-host  
(have the same source/destination host addresses)
  - ▶ process-to-process  
(have the same source/destination host/port pairs)
- ▶ you might hear the later situation called a **channel**
- ▶ one last thing: flows can be **explicit** or **implicit**
  - ▶ implicit flows are set up ad hoc by routers who just happen to notice a sequence of packets from one host to another
  - ▶ explicit flows are set up, but don't provide the reliability or delivery of a proper connection

# Fair Resource Allocation

- ▶ a good first step towards controlling congestion would be allocating network resources fairly
- ▶ what is “fair”?
  - ▶ a reservation-based resource allocation scheme provides an explicit way to create unfairness
  - ▶ with such a scheme, we might use reservations to enable a video stream to receive 1 Mbps across some link while a file transfer receives only 10 Kbps over the same link
  - ▶ probably not fair, even if it might be effective

# Jain's Fairness Index

- ▶ Raj Jain<sup>1</sup>, proposed a new fairness index to quantify the fairness of a resource allocation scheme:
  - ▶ given a set of flows  $1, 2, \dots, n$  with throughputs  $x_1, x_2, \dots, x_n$  (measured in consistent units), Jain's Fairness Index assigns the following index

$$f(x_1, x_2, \dots, x_n) = \frac{(\sum_{i=1}^n x_i)^2}{n \sum_{i=1}^n x_i^2}$$

- ▶ the index always results in a value between 0 and 1, with 1 representing perfect fairness

---

<sup>1</sup>Professor at Washington University with excellent notes online for networks and network security: <http://www.cse.wustl.edu/~jain/>