

Calculating Delays

Section 1.4

Delay, Loss, and Throughput in Packet-Switched Networks


Module Goals

At the conclusion of this module, students will be able to

- ▶ explain the difference between computing values that rely on both powers of 2 and powers of 10
- ▶ explain the sources of delay in network communications
- ▶ calculate delays in various circumstances

Performance

- ▶ the predominant metric used to characterize a network's performance is **bandwidth**
 - ▶ not a range of frequencies!
 - ▶ the number of bits per second that can be transmitted over a communication link
- ▶ $1 \text{ Mbps} = 10^6 \text{ bits/second} \approx 2^{10} \text{ bits/second}$
- ▶ 10^{-6} seconds to transmit each bit
(each bit occupies a $1 \mu\text{s}$ space on the wire)



Traditional Prefixes

- ▶ working with binary objects, we stole perfectly well-defined prefixes and redefined them using powers of 2 instead of powers of 10:
 - ▶ Kilo: traditionally 10^3 , now $2^{10} = 1024$
 - ▶ Mega: traditionally 10^6 , now $2^{20} = 1048576$
 - ▶ Giga: traditionally 10^9 , now $2^{30} = 1073741824$
- ▶ using the same prefix for approximately equivalent relationship

Enter Binary Prefixes

- ▶ in 1998, the International Electrotechnical Commission (IEC) standardized a suite of units to eliminate this confusion
- ▶ the power of 2 prefixes are given new names based on a portmanteau of the SI prefix and the “bi” from binary:
 - ▶ Kibi: 2^{10}
 - ▶ Mebi: 2^{20}
 - ▶ Gibi: 2^{30}
- ▶ and the SI prefixes are all powers of 10 again
- ▶ **network people always work with the power of 10 prefixes!**

Store and Forward

- ▶ packet-switched devices use **store-and-forward** processing
- ▶ the device does not receive a bit and immediately retransmit it; instead, every bit of the packet must arrive (and be stored) and only then does it transmit (or forward) the packet
- ▶ the device needs to analyze the packet's content to process it
- ▶ introduces some delay to the process

Calculating Delay

There are four delays experienced by a packet:

- ▶ **propagation delay**: signals travel extremely fast over links, but signals can not travel infinitely fast
- ▶ **transmission delay**: bits can only be pushed onto a link at a certain rate, limited by the power of the hardware
- ▶ **queuing delay**: packet switches along the way may incur delay because of packets in the outgoing queue
- ▶ **processing delay**: packet switches take some time to determine to which of the outgoing interfaces to route the packet

Calculating Delay

Propagation Delay

- ▶ once a bit is pushed into the link, it needs time to propagate to the next network device
- ▶ the **propagation time** t_{prop} is simply the distance between the source and destination divided by the **propagation rate** (distance over time)
(easy to remember with dimensional analysis: $d/d/s = s$)
- ▶ the propagation rate is dependent on the medium:
 - ▶ $\approx 2.3 \cdot 10^8$ m/s for copper
 - ▶ $\approx 3.0 \cdot 10^8$ m/s for wireless

Calculating Delay

Transmission Delay

- ▶ networking hardware can only push bits into a link at a limited rate, or **bandwidth** (its not a truck, but a series of tubes)
(802.11n @ 600 Mbit/s, Gigabit Ethernet @ 1 Gbit/s, USB3 @ 5 Gbit/s)
- ▶ the **transmission time** t_{xmit} is the amount of time it takes to transmit the bits of a packet onto the link
- ▶ if a packet contains L bits of data and is transmitted over a link with a bandwidth R (in bits per second), it takes L/R seconds to transmit the packet
- ▶ here, “transmit” only means pushing the packet onto the link; the packet must still propagate to the destination!

Calculating Delay

Queuing Delay

- ▶ generally, packets are transmitted in the order that they arrive
- ▶ while waiting to be transmitted, the packet is queued and experiences delay
 - ▶ if the queue is empty, the packet moves directly to transmission and experiences no queuing delay
 - ▶ if traffic is heavy, the queuing delay may be significant
 - ▶ if traffic is too heavy, the queue may fill and a packet might actually be lost
- ▶ queuing time is hard to quantify since it is so traffic dependent, so we'll just call it t_{queue} and specify it when/if it matters

Calculating Delay

Processing Delay

- ▶ typically, packets must be introspected to determine where to direct the packet, if there were any bit-level errors that occurred during a previous transmission, etc.
- ▶ this **processing time** t_{proc} depends strongly on the protocol so we'll ignore it for now
- ▶ processing delays on high-speed routers are usually strongly bounded, and are on the order of microseconds

Calculating Delay

The total delay for a single node is then

$$t_{nodal} = t_{proc} + t_{queue} + t_{xmit} + t_{prop}$$


Some implications:

- ▶ small transmissions \rightarrow propagation is important
- ▶ large transmissions \rightarrow bandwidth is important

Store and Forward Delay

Let's assume:

- ▶ there is just one device between the source and destination
- ▶ there is just one packet to send
- ▶ processing, queuing, and propagation delays are negligible



It takes L/R s to send the complete packet to the switch, then L/R s to set the complete packet to the destination.

Store and Forward Delay

What if there are multiple packets to send?

1. (1) sent to the switch (L/R s)



2. (2) sent to the switch, (1) sent to the destination (L/R s)



3. (2) sent to the destination (L/R s)




Store and Forward Delay

- ▶ what is the end-to-end delay for sending a single packet over a path with N links?
- ▶ what is the end-to-end delay for sending a P packets over a path with N links?

Store and Forward Delay

What if links have different transmission rates (bandwidths)?



- ▶ $R_1 < R_2$: the second link is limited by the rate of the first link
- ▶ $R_1 > R_2$: the second link will get overwhelmed by the first link
- ▶ end-to-end transmission rate becomes $\min(R_1, R_2)$
(assuming a lot of things, e.g., queues are infinitely big if $R_1 > R_2$)

Round-Trip Time (RTT)

- ▶ the previous equation was for one-way latency/delay
- ▶ most of the time there is a paired request/response—the latency there and back is also important

```
rover-224-252:~ jshiebel$ ping google.com
PING google.com (172.217.5.14): 56 data bytes
64 bytes from 172.217.5.14: icmp_seq=0 ttl=53 time=13.960 ms
64 bytes from 172.217.5.14: icmp_seq=1 ttl=53 time=14.478 ms
64 bytes from 172.217.5.14: icmp_seq=2 ttl=53 time=15.696 ms
64 bytes from 172.217.5.14: icmp_seq=3 ttl=53 time=15.210 ms
64 bytes from 172.217.5.14: icmp_seq=4 ttl=53 time=16.050 ms
^C
—— google.com ping statistics ——
5 packets transmitted, 5 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 13.960/15.079/16.050/0.769 ms
```

Link Utilization

Why pay for capability you don't use?

- ▶ we should try to use all of the carrying capacity that a link offers
- ▶ if I send data and then wait for a response before continuing, I'm wasting significant capacity!
- ▶ instead, we “ask forgiveness rather than permission”
(good motto for networks, bad motto for life)
- ▶ fill the pipe and deal with problems as they come

Link Utilization

- ▶ the carrying capacity of a particular link is how much data can be “in flight” between the source and destination
- ▶ how much is that?
 - ▶ if you want to know how many miles you’ve driven, you multiply the speed of the car by the time you’ve been driving
 - ▶ this is the produce of the bandwidth and propegation delay (push bits as fast as you can until the link is “full”)
- ▶ but remember that we send data and then are waiting for a response
- ▶ that’s two one-way delays, or the round-trip time

Link Utilization

- ▶ how much data is “buffered” for a 100 km, 10 Mbps link?
 - ▶ propagation delay: $100 \text{ km} / 2.3 \cdot 10^8 \text{ m/s} = 0.00043478 \text{ s}$
(0.00086957 s RTT)
 - ▶ capacity: $0.00086957 \text{ s} \cdot 10000000 \text{ bps} = 8696 \text{ bits}$
(1087 bytes, or **octets**)
- ▶ not a particularly interesting value by itself
- ▶ we'll learn later about the “sliding window protocol” that is heavily reliant on this number
- ▶ the capacity, combined with the size of the packet, tells us how many packets we will set before we expect a response from the receiver for the first packet