

# The Internet Protocol (IP)

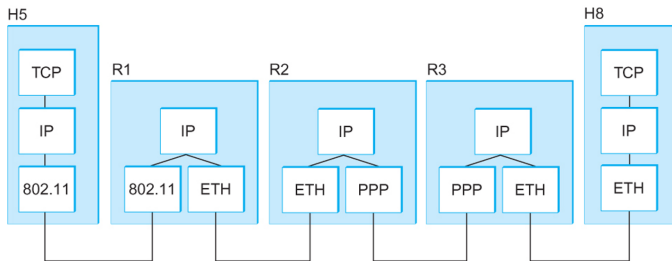
# Module Goals

At the conclusion of this module, students will be able to

- ▶ describe the IP addressing scheme
- ▶ describe how IP addresses are mapped to hardware addresses
- ▶ explain how IP addresses are dynamically assigned

# The Internet Protocol

- ▶ the key network protocol is simply the Internet Protocol (IP)
- ▶ key tool used today to build scalable, heterogeneous networks
- ▶ it runs on all nodes in a collection of networks and defines the infrastructure that allows these nodes and networks to function as a single, logical internetwork



# IP Service Model

- ▶ a “low expectations”, least common denominator kind of service (runs over anything)
- ▶ packet delivery model:
  - ▶ connectionless model for data delivery
  - ▶ best-effort delivery (unreliable service)
    - ▶ packets are lost
    - ▶ packets are delivered out of order
    - ▶ duplicate copies of a packet are delivered
    - ▶ packets can be delayed for a long time
- ▶ global addressing scheme (we'll talk about this later)

# Packet Format

0	3	4	7	8	15	16	18	19	31
Version	Header Length		TOS			Length			
Ident					Flags	Offset			
TTL		Protocol			Checksum				
Source Address									
Destination Address									
Options (If Any)									
Data Octets									


# Packet Format

- ▶ Version (4 bits): currently 4
- ▶ Header Length (4 bits): length of the header in 32-bit words
- ▶ TOS (8 bits): type of service, now the differentiated services field
- ▶ Length (16 bits): number of bytes in this datagram
- ▶ Ident/Flags/Offset (32 bits): used by fragmentation
- ▶ TTL (8 bits): number of hops this datagram has traveled
- ▶ Protocol (8 bits): demux key; TCP is 6, UDP is 17
- ▶ Checksum (16 bits): header checksum only
- ▶ Destination Address (32 bits): destination address
- ▶ Source Address (32 bits): source address

# Fragmentation and Reassembly

- ▶ each link layer has some maximum transmission unit (MTU)  
(Ethernet: 1500 bytes; FDDI: 4500 bytes)
- ▶ if IP is going to run on any of these link layers (and has to be contained in the payload of link layer frames, we have to figure out how to size packets
  - ▶ option 1: make them sufficiently small so that we hope they fit in anything
  - ▶ option 2: allow packets to get **fragmented** between multiple frames and reassembled at the destination
    - ▶ all fragments carry the same identifier in the Ident field
    - ▶ fragments are self-contained datagrams
    - ▶ if one fragment gets lost, we discard all fragments
    - ▶ we do not try to recover anything from missing fragments

# Example Fragmentation



Host H5 wants to send a datagram to Host H8 through the network.



# IP Fragmentation and Reassembly

- ▶ the original datagram is in (a) (a)
- ▶ when it reaches router R2, it gets fragmented into 3 datagrams (b)
  - ▶ Ident is set to some value  $x$
  - ▶ Flags “more bit” is set
  - ▶ Offset is set for the number as the starting point in terms of 8 byte chunks (fragmentation in IP is always on 8 byte boundaries)

Start of Header			
$x$		0	0
Rest of Header			
1400 Data Bytes			

Start of Header			
$x$		1	0
Rest of Header			
512 Data Bytes			

Start of Header			
$x$		1	64
Rest of Header			
512 Data Bytes			


Start of Header			
$x$		0	128
Rest of Header			
376 Data Bytes			

# Global Addresses

- ▶ what makes a good addressing scheme?
- ▶ step 1: globally unique addresses
  - ▶ TCP had ports, but these weren't unique in the world
  - ▶ we'll discuss hardware addresses for protocols like Ethernet which do have globally unique addresses, but...
  - ▶ these addresses don't provide any meaningful information on how to actually find these devices
- ▶ what if the addresses provided more information?

# Hierarchical Addresses

- ▶ hierarchical approach: network + host
- ▶  $\approx 4$  billion addresses (32 bits)



- ▶ specified using dot notation:  
10.3.2.4, 128.96.33.81, 192.12.69.77
- ▶ most addresses are now classless because the original class scheme proved to be too inflexible

# A Flawed Idea

- ▶ it quickly became clear that the network/host division of IP addresses wasn't sufficient
- ▶ a network with 6 hosts would need a class C address, which supports up to 255 hosts (wasting 249! addresses)
- ▶ if it was a class B address we would be wasting 64,000 addresses
- ▶ in a nutshell, assigning whole blocks of network addresses to a real, physical network is inefficient

# Working Towards a Solution

- ▶ one solution: **subnets**  
(but these only fix part of the problem)
  - ▶ subnetting lets us take a classful address among multiple smaller networks
  - ▶ example: a given organization might not need a class B address (for 64K hosts) and gets 16 class C addresses instead (4096 hosts)
  - ▶ now we have a different problem: the 16 different address ranges need 16 different forwarding entries in every router (tradeoff between forwarding entries and address waste)
- ▶ a better solution: **Classless InterDomain Routing (CIDR)**

# Classless Addressing

- ▶ CIDR tries to balance the desire to minimize the number of routes that a router needs to know against the need to hand out addresses efficiently
- ▶ CIDR uses **aggregate routes**
  - ▶ use a single entry in the forwarding table to tell the router how to reach a lot of different networks
  - ▶ breaks the rigid boundaries between address classes
  - ▶ essentially lets any number of bits in the address represent the “network” portion of the address

# Classless Addressing

- ▶ consider that hypothetical org needing those 16 class C network ranges
- ▶ instead of handing out 16 addresses at random, hand out a block of contiguous class C addresses
- ▶ suppose we assign the class C network numbers from 192.6.16 through 192.4.31
  - ▶ if we blow those values up into binary, we see the top 20 bits of all the addresses in this range are the same: 11000000  
00000100 0001
  - ▶ in essence, we just created a 20 bit network number, which is in between a class B network number and class C number
- ▶ of course, this means somebody needs to hand out blocks of class C addresses to organizations such that they share a


# CIDR Notation

- ▶ this calls for a slightly different way of representing the address
- ▶ the convention is to place a /X after the prefix where X is the prefix length in bits
- ▶ for example, the 20 bit prefix for all the networks 192.4.16 through 192.4.31 is represented as 192.4.16/20
- ▶ by contrast, if we wanted to represent a single class C network number, which is 24 bits long, we would write it as 192.4.16/24
- ▶ Michigan Tech owns 141.219/16



# Classless Addressing

- ▶ this can be done hierarchically too
- ▶ the ISP can be allocated the IPs and then hand them out to customers in smaller blocks
- ▶ the greater internet sees 128.112.128/21
- ▶ For each of the clients, the ISP knows 128.112.xxx/yy  
(yy doesn't need to be the same for all the customers serviced by an ISP)



# Multiple Matches

- ▶ as we shared, Michigan Tech owns 141.219/16
- ▶ however, RIPE Network Coordination Center owns 141/8
- ▶ should 141.219.10.20 be routed to RIPE or MTU?
  - ▶ principle of **longest specific match**
  - ▶ the datagram should be routed to MTU
  - ▶ you can look up routing info using `route` on Unix systems

# Delivering a Datagram

- ▶ so every device has an IP address
- ▶ a host on some network sending a datagram to a host on another network can get that datagram to the network
- ▶ if it's not on my network, I send it to the default router—the **gateway**
- ▶ then it's up to other devices to get it to the gateway for the destination network
- ▶ but how does it get to the actual host?
  - ▶ an Ethernet switch doesn't know about IP addresses (switches are data link layer devices!)
  - ▶ somehow we have to translate an IP address into an address that “makes sense” on the network the datagram is traveling on

# Address Translation

- ▶ idea 1: encode the physical address (e.g., Ethernet device address) into the IP address
  - ▶ question: how do you encode the 48 bit Ethernet address into a 32 bit IP address?
  - ▶ answer: you don't; come up with a better ideaa
- ▶ idea 2: address translation table
  - ▶ a host maintains a table mapping IP addresses to physical addresses
  - ▶ question: how does the table get populated?
  - ▶ answer: either a human does it, or the hosts dynamically learn the network
  - ▶ the dynamic version is the **Address Resolution Protocol (ARP)**

# Address Resolution Protocol

- ▶ imagine: Host A (10.0.1.4) wants to send a datagram to Host B (10.0.1.8)
- ▶ it looks in the table and has no match for that IP
- ▶ it sends out a broadcast message on the network: “Who is 10.0.1.8?”
- ▶ reply from Host B: “10.0.1.8” is at address 68:A8:6D:54:D8:CA
  - ▶ Host A adds that information to its ARP table
  - ▶ Host B also just got a frame from Host A with its IP and physical address, which B adds to its own ARP table (after all, it's pretty likely there's future communication in store)
- ▶ what about the other hosts watching the ARPing?  
(they generally ignore it)

# Packet Format

0	7	8	15	16	31
Hardware Type			Protocol Type		
H. Length		P. Length		Operation	
Source Hardware Address					
Source Hardware Address			Source Protocol Address		
Source Protocol Address			Target Hardware Address		
Target Hardware Address					
Target Protocol Address					

# Packet Format

- ▶ Hardware Type (16 bits): type of physical network (e.g., Ethernet = 1)
- ▶ Protocol Type (16 bits): type of higher layer protocol (e.g., IP = 0x0800)
- ▶ H. Length: length of hardware addresses, in bytes (e.g., Ethernet = 6)
- ▶ P. Length: length of protocol addresses, in bytes (e.g., IPv4 = 4)
- ▶ Operation: request (1), reply (2)
- ▶ Sender Hardware Address: media address of the sender
- ▶ Sender Protocol Address: internetwork address of the sender
- ▶ Target Hardware Address: media address of the target
- ▶ Target Protocol Address: internetwork address of the target

# ARP Security

- ▶ does ARP seem secure?
- ▶ there are definitely issues
- ▶ what if a malicious host responds, “sure, I’ve got IP x.x.x.x”, even when it doesn’t?
- ▶ there is also **gratuitous ARP** that is sent by a device first joining the network
- ▶ what if a malicious host sends a **gratuitous ARP** with the IP of the gateway?



# Getting IP Addresses

- ▶ how do devices get IP addresses anyways?
- ▶ MAC addresses are given to devices by the manufacturer  
(could IP addresses be given out this way?)
- ▶ two main ways this is handled: static and dynamic IP addresses

# Dynamic Host Configuration Protocol (DHCP)

- ▶ the DHCP server is responsible for providing configuration information to hosts
- ▶ there is at least one DHCP server for an administrative domain
- ▶ DHCP server maintains a pool of available addresses

# DHCP Requests

- ▶ newly booted or attached hosts send a DHCPDISCOVER message to a special IP address (255.255.255.255)
- ▶ routers refuse to forward frames containing this IP address to external networks
- ▶ DHCP relay agent unicasts the message to the DHCP server and waits for the response
- ▶ why a relay?
- ▶ we have a bunch of small networks hooked together via bridges, and who wants a DHCP server for every little network?
- ▶ addresses are leased for a certain amount of time, though a host can renew its lease

# DHCP Responses

The DHCP Server tells the host:

- ▶ what IP address to use
- ▶ how long we can use the IP address for
- ▶ local DNS servers
- ▶ domain name information
- ▶ subnet mask
- ▶ default gateway

# DHCP Security

- ▶ does DHCP seem secure?
- ▶ we could have an unauthorized DHCP server sending out information ranging from useless to malicious
- ▶ with a little work, we could exhaust the IP pool by sending loads of DHCP requests

# Summary

- ▶ the network layer might have to fragment packets to fit within a link layer transmission
- ▶ the Internet Protocol requires globally unique addresses
- ▶ addresses can be classful vs. classless
- ▶ address discovery can be done with ARP
- ▶ address allocation can be done with DHCP
- ▶ if everybody can see your traffic, and if you rely on other devices for correctness, you have to show a lot of trust